



NANYANG TECHNOLOGICAL UNIVERSITY

EE6222 Machine Vision

Assignment 1: Dimensionality Reduction for Classification

MSc in Electrical and Electronic Engineering

Yan Ziming

G2507084J

October 23, 2025

Contents

1	Introduction	2
2	Dataset and Preprocessing	2
2.1	Dtaset Description	2
2.2	Normalization and Standarlization	2
2.3	Dataset Partition	3
3	Methodology	3
3.1	Linear Dimensionality Reduction	3
3.1.1	Linear Discriminant Analysis	3
3.1.2	Principal Component Analysis	5
3.1.3	Fisherfaces	6
3.2	Nonlinear Dimensionality Reduction	6
3.2.1	Locally Linear Embedding	6
3.2.2	Isometric Mapping	8
3.2.3	Pairwise Controlled Manifold Approximation	9
3.3	Classifiers Establishment	10
3.3.1	Minimun Mahalanobis Distance Classifier	10
3.3.2	K-Nearest Neighbors Classifier	11
3.3.3	Logistic Regression Classifier	11
3.4	Evaluation Methods	11
4	Results and Discussion	11
5	Conclusion	11

1 Introduction

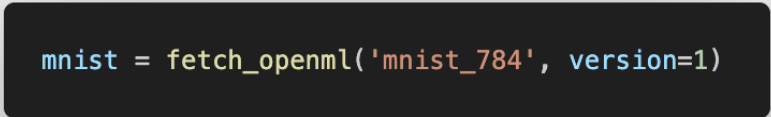
In this Assignment, I selected the MINIST dataset for training and testing the performance of classifiers. As for the dimension of images are large, I explored several dimensionality reduction methods to reduce the dimension of data while trying to retain the most important features, including **PCA**, **LDA**, combined method (Fisherfaces), **LLE**, **Isomap**, and **PaCMAP**. I compared the effectiveness of these methods by plotting their performance with varying number of dimensions.

In addition, I establish several basic classifiers, like **MMD**, **KNN**, and include Logistic Regression Classifier with sklearn package. Therefore, I not only compare the effectiveness of different dimensionality reduction methods, but also access the accuracy of classify images for these classifier methods. I would describe the details of my experiment in the following sections.

2 Dataset and Preprocessing

2.1 Dataset Description

The MINIST dataset contains **70,000** images of handwritten digits (0-9) with a resolution of 28x28 pixels, which means there are **10** classes intotal. Each image had been saved with the format of vector with **784 in length**. In addition, the images are only have one channel, which means they are grayscale pixels. Because of the pre-shaped vectors, single channel, and easy to get features, I think MINIST dataset is suitable for this assignment. Therefore, I fetch the dataset from sklearn package directly 1.



```
mnist = fetch_openml('mnist_784', version=1)
```

Figure 1: Fetch Code

2.2 Normalization and Standarlization

Since the MINIST dataset had been well preprocessed and tested by plenty of researchers, I didn't do too much basic preprocessing, like outlier and missing value detection. However, I though normalization and standarlization arer still neccessary for the following basic classifiers, which may be affected intensely by the scle of features. So I rearrange the vectors to 0 mean and unit variance.

$$\bar{X} = \frac{X}{L-1}$$
$$X_{std} = \frac{X - \mu}{\sigma}$$

2.3 Dataset Partition

I found the number of samples (70,000) is too large, which may lead to many difficulties in my following experiment. Firstly, I will training several classifiers with many dimension reduction methods, meaning the number of training process will be big. If I train 50,000-70,000 for each training, the consumed time will be very long. Secondly, my classifiers which will be used are simple and basic, therefore, they are easy to converge with small scale of data. With these two reason, I decide to extract 7,000 samples from the origin dataset formly to reduce the expence of time.

After that, I split the dataset into training set and testing set through *train_test_split* function with ratio of 6:1 and finished the process of preparation of dataset 2.

```
Training Feature Dimension: (6000, 784)
Traing Label Dimension: (6000,)
Testing Feature Dimension: (1000, 784)
Testing Label Dimension: (1000,)
```

Figure 2: Dimension of practical dataset

3 Methodology

3.1 Linear Dimensionality Reduction

3.1.1 Linear Discriminant Analysis

Linear Discriminant Analysis (LDA) aims to find a linear combination of features with weight matrix $\mathbf{w}$ to transform the data into linear space, ensuring that the separation between different classes is maximized while the variance within each class is minimized.

Step 1: Mean Value Computation

With given m classes $\{w_1, w_2, \dots, w_m\}$, there are c samples in each class w : $\{X_1, X_2, \dots, X_c\}$. For each sample vector, it has n features, $\{x_1, x_2, \dots, x_n\}$. In this experiment, $n=784$.

Overall Mean:

$$\mu_t = \frac{1}{N} \sum_{i=1}^N x_i$$

Class Mean:

$$\mu_i = \frac{1}{n_i} \sum_{j=1}^{n_i} x_j$$

Where N is the total number of samples, and n_i is the number of samples in class w_i .

Step 2: Scatter Matrix Computation

Between-Class Scatter Matrix:

$$S_B = n_i * \sum_{j=1}^m (\mu_j - \mu_t)(\mu_j - \mu_t)^T$$

Within-Class Scatter Matrix:

$$S_W = \sum_{j=1}^m \sum_{x_i \in w_j} (x_i - \mu_j)(x_i - \mu_j)^T$$

```
Sw = np.zeros((n_features, n_features))
for i, c in enumerate(self.classes):
    Xc = X[y == c] - means[i]
    Sw += Xc.T @ Xc
Sw += 1e-6 * np.eye(n_features) #invertable
Sb = np.zeros((n_features, n_features))
for i, c in enumerate(self.classes):
    n_c = (y == c).sum()
    mean_diff = means[i] - self.overall_mean.reshape(-1, 1)
    Sb += n_c * (mean_diff @ mean_diff.T)
```

Figure 3: Scatter Computation Code

Step 3: Weight Matrix Computation

As I mentioned before, the LDA aims to find the optimal weight matrix to separate classes on two dimension, equaling to the problem: maximize the ratio of between-class scatter to within-class scatter.

$$\mathbf{w}^* = \arg \max_{\mathbf{w}} \frac{|\mathbf{w}^T S_B \mathbf{w}|}{|\mathbf{w}^T S_W \mathbf{w}|} \quad (1)$$

Since the $S_B \mathbf{w}_i = \lambda_i S_W \mathbf{w}_i$,

$$\lambda \mathbf{w} = S_W^{-1} S_B \mathbf{w} \quad (2)$$

Therefore, the $\mathbf{w}$ is the eigenvectors matrix according to eigenvalue λ , the largest λ refers to best matrix $\mathbf{w}^*$.

```
eigvals, eigvecs = eigh(Sb, Sw)
sorted_idx = np.argsort(eigvals)[::-1]
max_comp = min(len(self.classes) - 1, n_features)
n_comp = max_comp if self.n_components is None else min(self.
n_components, max_comp)
self.scalings = eigvecs[:, sorted_idx[:n_comp]]
```

Figure 4: Eigenvector Matrix Computation

Step 4: Transform Dataset

After getting the weight matrix, the transformation of input data $\mathbf{X}$ is:

$$\mathbf{Y} = \mathbf{X}\mathbf{w}^*$$

In practice, we will choose the k^{th} value of sorted eigenvalues to form $\mathbf{w}^*$. Theoretically, LDA is designed to project samples onto a lower-dimensional subspace, with the dimensionality being at most $C - 1$, because the rank of S_B is at most $C-1$, C is the number of classes.

3.1.2 Principal Component Analysis

Similar to LDA, Principal Component Analysis (PCA) is another effective method to dimensionality reduction. The difference is PCA can project the input data to different dimensions without limitation like LDA. The projected data after PCA will also separate largely, which means the redundancy elements are eliminated.

Step 1: Data Processing

Before the projection, the data should be scaled and normalized, eliminating bias. Specifically, compute mean value to shift to 0 mean and covariance to investigate the distribution situation of data.

Covariance Matrix Computation:

$$\Sigma_i = \sum_{x_j \in w_i}^{n_i} (x_j - \mu_i)(x_j - \mu_i)^T$$

Where μ_i is the mean value of class w_i

Step 2: Eigenvectors Computation

The PCA aims to find the direction $\mathbf{w}$, the variance of data projected to that direction is largest. For the transformed data $Y = \mathbf{w}X$, variance computed:

$$Var(Y) = \frac{1}{N} \sum_{i=1}^N (w^T x_i)^2 = w^T \Sigma w$$

Solving the $\arg \max_w (w^T \Sigma w)$,

$$\Sigma \mathbf{w} = \lambda \mathbf{w} \tag{3}$$

Larger eigenvalue refers to better eigenvector, by sort the eigenvalues in descending sequence, we can fetch k^{th} large vlues and their correspondingly eigenvectors to form best matrix $\mathbf{w}^*$.

```

self.mean = np.mean(X, axis=0)
X_centered = X - self.mean
covariance_matrix = np.cov(X_centered, rowvar=False)
eigenvalues, eigenvectors = np.linalg.eigh(covariance_matrix)
sorted_indices = np.argsort(eigenvalues)[::-1]
self.components = eigenvectors[:, sorted_indices[:self.
n_components]]

```

Figure 5: PCA Realization Code

Step 3: Transform Dataset

With above process, we can get the corresponding optimal weight matrix, according to the input data and componenet coefficient k . Applying the matrix to input data, transformation finished.

3.1.3 Fisherfaces

Based on learning in course, I acknowledge that the reduction of feature equals to deleting redundant features, which can increase the accuray of models. However, when too much features deleted, the remaining features may not includes all the information we need, which causes the decrease of accuray backward. Therefore, the best accuracy of LDA towards high dimension data is much lower than PCA normally. To overcome this situation, the Fisherfaces, which combines the LDA and PCA method is developed.

For realization, I just handle the data with PCA to roughly decrease the dimension firstly. Then, I applied LDA based on the corresponding PCA space to finish the whole reduction.

3.2 Nonlinear Dimensionality Reduction

A **manifold** is a topological space that locally resembles Euclidean space. Intuitively, even if data lie in a high-dimensional ambient space $\mathbb{R}^D$, their intrinsic degrees of freedom may be much smaller, forming a d -dimensional manifold ($d \ll D$). For example, images of a rotating object under fixed lighting conditions can be considered points on a 2D manifold embedded in a high-dimensional pixel space.

Formally, a dataset $X = \{x_1, x_2, \dots, x_N\}$ is assumed to lie on or near a smooth manifold $\mathcal{M}$ of dimension d . The goal of **manifold learning** is to find a mapping $f: \mathbb{R}^D \rightarrow \mathbb{R}^d$ that preserves intrinsic geometric relationships among points.

Unlike linear methods such as PCA, which only capture global linear structures, manifold learning methods exploit the underlying nonlinear geometry. I include manifold learning methods for dimension reduction in this experiment.

3.2.1 Locally Linear Embedding

Brief Introduction:

Locally Linear Embedding (LLE) is a nonlinear dimensionality reduction algorithm that preserves local geometric relationships between neighboring data points. It assumes that each data point and its neighbors lie on or close to a locally linear manifold. LLE first reconstructs each data point as a linear combination of its nearest neighbors in the original high-dimensional space, then seeks a low-dimensional embedding that maintains these same reconstruction weights [1].

Let the dataset be $\{x_1, x_2, \dots, x_N\}$, where each $x_i \in \mathbb{R}^D$. For each data point x_i , find its K nearest neighbors $\mathcal{N}(i)$ using Euclidean distance. Each point is reconstructed from its neighbors:

$$x_i \approx \sum_{j \in \mathcal{N}(i)} w_{ij} x_j$$

The reconstruction weights w_{ij} are obtained by minimizing the local reconstruction error:

$$\min_W \sum_{i=1}^N \left\| x_i - \sum_{j \in \mathcal{N}(i)} w_{ij} x_j \right\|^2 \quad (4)$$

subject to the constraint:

$$\sum_{j \in \mathcal{N}(i)} w_{ij} = 1$$

These weights capture the intrinsic geometry of the data.

In the embedding step, we find low-dimensional representations $y_i \in \mathbb{R}^d$ ($d \ll D$) that preserve the same reconstruction weights:

$$\min_Y \sum_{i=1}^N \left\| y_i - \sum_{j \in \mathcal{N}(i)} w_{ij} y_j \right\|^2 \quad (5)$$

subject to:

$$\sum_i y_i = 0, \quad \frac{1}{N} \sum_i y_i y_i^T = I$$

The optimal embedding $Y = [y_1, \dots, y_N]^T$ is obtained by computing the bottom $(d+1)$ eigenvectors of the matrix:

$$M = (I - W)^T (I - W) \quad (6)$$

and discarding the smallest eigenvector corresponding to the zero eigenvalue.

Implementation:

In this experiment, I implemten the LLE by using the `LocallyLinearEmbedding` class from `scikit-learn`. The key parameters are the number of neighbors (`n_neighbors`) and the target embedding dimension (`n_components`). Below shows an example code used in this experiment:


```

reduction_m = LocallyLinearEmbedding(n_components=n)
reduction_m.fit(X_train)
X_train_reduced = reduction_m.transform(X_train)
X_test_reduced = reduction_m.transform(X_test)

```

Figure 6: LLE Realization Code

3.2.2 Isometric Mapping

Brief Introduction:

Isometric Mapping (Isomap) is a nonlinear dimensionality reduction method based on the manifold learning assumption, which posits that high-dimensional data lie on a low-dimensional nonlinear manifold. Unlike linear methods such as PCA, Isomap seeks to preserve the intrinsic geometry of the manifold by approximating geodesic distances between points [2].

For a dataset $\{x_1, x_2, \dots, x_N\} \subset \mathbb{R}^D$, Isomap constructs a neighborhood graph connecting each point to its K nearest neighbors. The edges are weighted by the Euclidean distance between connected points. The geodesic distance between any two points x_i and x_j is then approximated by the shortest path distance in this graph:

$$D_G(i, j) = \min_{\text{paths } i \rightarrow j} \sum_{(p, q) \in \text{path}} \|x_p - x_q\| \quad (7)$$

Once the matrix of geodesic distances D_G is obtained, classical Multidimensional Scaling (MDS) is applied to compute a low-dimensional embedding $Y \in \mathbb{R}^{N \times d}$ that best preserves these distances. The centered inner-product matrix is calculated as

$$B = -\frac{1}{2}HD_G^2H, \quad \text{where } H = I - \frac{1}{N}\mathbf{1}\mathbf{1}^T \quad (8)$$

and the embedding is obtained from the top d eigenvectors of B :

$$Y = V_d \Lambda_d^{1/2}$$

where Λ_d contains the d largest eigenvalues, and V_d the corresponding eigenvectors. The resulting low-dimensional coordinates preserve the manifold's intrinsic global geometry.

Implementation:

Isomap can be implemented using the Isomap class from `scikit-learn`. The main parameters are the number of neighbors (`n_neighbors`) and the target embedding dimension (`n_components`). The code snippet used in this work is shown in Figure 7, and the resulting 2D embedding is visualized for comparison with other methods.

```

reduction_m = Isomap(n_components=n)
reduction_m.fit(X_train)
X_train_reduced = reduction_m.transform(X_train)
X_test_reduced = reduction_m.transform(X_test)

```

Figure 7: Isomap Realization Code

3.2.3 Pairwise Controlled Manifold Approximation

Brief Introduction:

PaCMAP is a modern nonlinear dimensionality reduction method designed to preserve both local and global structures in high-dimensional data. Unlike traditional methods such as t-SNE and UMAP, which primarily focus on local structure, or TriMAP, which emphasizes global relationships, PaCMAP balances these by controlling three types of point-pair relationships:

- **Neighbor pairs:** Preserve local neighborhood structure.
- **Mid-near pairs:** Preserve medium-range relationships.
- **Far pairs:** Prevent collapse of distant points and maintain global separation.

The main objective of PaCMAP is to optimize a low-dimensional embedding such that these point-pair relationships are faithfully preserved, producing a high-quality representation suitable for visualization and downstream tasks [3]. PaCMAP’s objective function consists of three components corresponding to the three types of point pairs:

$$L_{\text{PaCMAP}} = L_{\text{neighbors}} + L_{\text{mid-near}} + L_{\text{far}} \quad (9)$$

where:

- $L_{\text{neighbors}}$ measures the discrepancy of neighbor pairs in low-dimensional space relative to the original space.
- $L_{\text{mid-near}}$ measures the preservation of mid-range distances.
- L_{far} ensures that distant points remain well separated.

The specific loss formulations and weighting schemes can be found in the original paper [3].

Implementation: As mentioned in to essay, developers have embeded PaCMAP in Python. We can implement the reduction method by using the pacmap package as follows:

```

embedding = pacmap.PaCMAP(
    n_components=n,
    n_neighbors=10,
    MN_ratio=0.5,
    FP_ratio=2.0
)
X_reduced=embedding.fit_transform(X)

```

Figure 8: PaCMAP Realization Code

Since it is a very new tool designed to reduce the dimension of data to visualize the relation between high dimension features, pacmap only allows to reduce the total dataset, different from the previous methods package usage, which can be seen by comparing the realization of PaCMAP (Figure 8) with Isomap (Figure 7) or LLE (Figure 6).

3.3 Classifiers Establishment

After reducing the dimension, I need to verify the influence of these reduction methods. Therefore, I need to develop classifiers and use their performance (accuracy) to judge the effect of datashape changes.

3.3.1 Minimum Mahalanobis Distance Classifier

Mahalanobis Distance represents the difference between two images. The MMD classifier refers to divide the input image data to the class which contains the nearest image data.

The Mahalanobis Distance of sample $X = \{x_1, x_2, \dots\}$ to class w_j :

$$D_j = (X - \mu_j) \Sigma_j^{-1} (X - \mu_j)^T$$

Where μ_j is the mean value of data in class, Σ_j is the covariance of class data.

There is possible that covariance may be singular, leading no inversion. So I added the covariance with $1e-6 * \text{np.eye}(X.\text{shape}[1])$ to avoid singular condition.

After getting the distance of input point to each classes, the predicted label of input point is considered to

$$\underset{w_j}{\operatorname{argmin}}(D(X, w_j))$$

3.3.2 K-Nearest Neighbors Classifier

3.3.3 Logistic Regression Classifier

3.4 Evaluation Methods

4 Results and Discussion

5 Conclusion

References

- [1] S. T. Roweis and L. K. Saul, “Nonlinear dimensionality reduction by locally linear embedding,” *Science*, vol. 290, no. 5500, pp. 2323–2326, 2000.
- [2] J. B. Tenenbaum, V. de Silva, and J. C. Langford, “A global geometric framework for nonlinear dimensionality reduction,” *Science*, vol. 290, no. 5500, pp. 2319–2323, 2000.
- [3] Y. Wang, Y. Sun, H. Huang, and C. Rudin, “Dimension Reduction with Locally Adjusted Graphs,” Dec. 2024. arXiv:2412.15426 [cs].

Appendix